

# Fair Distribution of Digital Payments: Balancing Transaction Flows for Regulatory Compliance

Ashlesha Hota\*  
IIT Kharagpur  
India

Abolfazl Asudeh  
University of Illinois Chicago  
United States

Shashwat Kumar\*  
IIT Kharagpur  
India

Palash Dey  
IIT Kharagpur  
India

Daman Deep Singh  
IIT Delhi  
India

Abhijnan Chakraborty  
IIT Kharagpur  
India

## Abstract

The concentration of digital payment transactions in just two UPI apps - PhonePe and Google Pay - has raised concerns of duopoly in India's digital financial ecosystem. To address this, the National Payments Corporation of India (NPCI) has mandated that no single UPI app should exceed 30% of total transaction volume. Enforcing this cap, however, poses a significant computational challenge: how to redistribute user transactions across apps without causing widespread user inconvenience while maintaining capacity limits?

In this paper, we formalize this problem as the MINIMUM EDGE ACTIVATION FLOW (MEAF) problem on a bipartite network of users and apps, where activating an edge corresponds to a new app installation. The objective is to ensure a feasible flow respecting app capacities while minimizing additional activations. We further prove that MINIMUM EDGE ACTIVATION FLOW is NP-complete. To address the computational challenge, we propose scalable heuristics, named DECOUPLED TWO-STAGE ALLOCATION STRATEGY (DTAS), that exploit flow structure and capacity reuse. Experiments on large semi-synthetic transaction network data show that DTAS finds solutions close to the optimal ILP within seconds, offering a fast and practical way to enforce transaction caps fairly and efficiently.

## Keywords

UPI transaction caps, Flow optimization, Fair distribution

## 1 Introduction

Digital payments in India have witnessed explosive growth through the Unified Payments Interface (UPI), an interoperable platform enabling seamless bank-to-bank transfers. UPI now facilitates 85% of all digital transactions in the country [9]. Despite the presence of more than 15 UPI applications, the market remains highly concentrated, with just two platforms PhonePe and Google Pay together accounting for over 80% of all transactions [7]. This concentration poses competition concerns and systemic risks, as an outage in a single app could disrupt a large share of India's digital payments.

To mitigate these risks, the National Payments Corporation of India (NPCI) introduced a market share cap requiring that no UPI app handle more than 30% of total transactions [4]. While this policy promotes fairness and innovation, its enforcement is nontrivial: user behavior drives transaction volumes, and users naturally prefer familiar apps. The central challenge is thus to regulate transaction flow across apps in real time without violating capacity limits or disrupting users.

\*Both authors contributed equally to this research.

A naïve enforcement mechanism is a tail-drop policy, blocking further transactions once an app exceeds its quota. Although effective, it leads to abrupt failures and poor user experience. A more practical solution is an alert-based strategy, where users are notified when an app nears its limit and encouraged to switch to another app. This idea parallels Random Early Detection (RED) in network congestion control [6], where early warnings prevent sharp throughput drops.

However, implementing this strategy faces user behavioral challenges: users are often reluctant to install and maintain multiple UPI apps, and simply providing alerts may not guarantee sufficient switching. Indian government already allocates budgetary incentives to support zero-fee UPI transactions. We propose that a portion of this subsidy be explicitly earmarked for encouraging installation and active usage of smaller UPI apps, for example through time-bound, usage-linked incentives for new users. Since such incentive budget would be limited, in this work, our focus is to minimize the additional app installations while achieving a balanced (regulation-compliant) transaction distribution and zero transaction failure.

We formalize this as the MINIMUM EDGE ACTIVATION FLOW problem on a bipartite flow network  $G = (V = \{s\} \cup U \cup A \cup \{t\}, E = E_{\text{solid}} \cup E_{\text{dashed}})$ . Here, users  $U$  and apps  $A$  form two partitions: solid edges represent existing app installations, dashed edges denote potential ones,  $s-U$  edges have capacities  $t_u$  (user transaction volumes), and  $A-t$  edges have capacities  $c_a$  (app limits). The objective is to activate the minimum number of dashed edges to ensure a feasible integral flow from  $s$  to  $t$  without exceeding any app's capacity.

Despite its intuitive formulation, the problem is computationally intractable: we prove that determining the smallest feasible activation set or equivalently, the minimal number of additional app installations is NP-complete, even when restricted to only three apps. This rules out any polynomial-time exact solution unless  $P = NP$ . Consequently, we propose efficient greedy heuristics that approximate the optimal activation pattern while being scalable for large transaction networks.

In summary, we make the following key contributions in this paper:

- (1) **Computational intractability.** We establish that the decision version of the problem is NP-complete via a polynomial-time reduction from 3-PARTITION.
- (2) **Greedy Heuristics.** To address computational intractability, we propose a linear programming (LP) relaxation of the integral flow constraints. The relaxation provides a tractable method to compute lower bounds and fractional flow solutions, which can guide edge activation decisions and inform practical redistribution strategies.

- 117 ▷ **CARL (CAPACITY-AWARE REUSE-FIRST LAYERED AL-**  
 118 **LOCATION):** The CARL heuristic prioritizes efficient  
 119 capacity utilization and reuse of already deployed ap-  
 120 plications. It operates by processing users in ascending  
 121 order of their transaction-to-capacity ratio, thereby giving  
 122 precedence to users with limited available capacity  
 123 relative to their transactional load. For each user, CARL  
 124 performs allocation in three structured layers: (i) prein-  
 125 stalled apps, (ii) existing extra apps, and (iii) new extra  
 126 apps if additional capacity is required. Within each  
 127 layer, apps are sorted by their remaining capacity in  
 128 descending order to ensure balanced load distribution.  
 129 This layered reuse-first design minimizes redundant app  
 130 installations, enhances overall system stability, and im-  
 131 proves scalability. By emphasizing high-capacity reuse  
 132 and adaptive layering, CARL achieves a judicious trade-  
 133 off between transaction coverage and capacity fairness  
 134 across users and applications.
- 135 ▷ **DTAS (DECOPLED TWO-STAGE ALLOCATION STRAT-  
 136 EGY):** The DTAS heuristic employs a two-phase greedy  
 137 framework that significantly improves upon CARL by  
 138 enforcing stricter separation between allocation tiers.  
 139 In *Phase 1*, all transactions are allocated exclusively  
 140 through preinstalled apps, ensuring that every user fully  
 141 utilizes their zero-cost connections before accessing  
 142 additional capacity. This prevents the premature allo-  
 143 cation of existing extra apps to a single user, which  
 144 in CARL could lead to early saturation of such apps  
 145 and consequently restrict access for other users who  
 146 have them preinstalled. In *Phase 2*, DTAS allocates the  
 147 remaining unmet demand by first reusing existing extra  
 148 apps (installed for other users) and then introducing  
 149 new extra apps only when necessary. By prioritizing  
 150 preinstalled capacity globally before activating new  
 151 edges, DTAS achieves superior capacity preservation,  
 152 fairness, and reduced installation overhead. Empiri-  
 153 cal evaluation shows that this disciplined two-phase  
 154 structure allows DTAS to outperform CARL in both  
 155 total feasible allocation and capacity balance across the  
 156 network.
- 157 (3) **Experimental Evaluation.** To empirically validate our pro-  
 158 posed allocation strategies, we conducted extensive exper-  
 159 iments across configurations ranging from 10,000 to 100  
 160 million transactions and up to 1.2 million users. Each setup  
 161 was tested under consistent capacity constraints with 20  
 162 varying transaction-to-user ratios. Both heuristics—CARL  
 163 (CAPACITY-AWARE REUSE-FIRST LAYERED ALLOCATION) and  
 164 DTAS (Dual-Tier Allocation Strategy)—were evaluated  
 165 against the optimal Integer Linear Program (ILP) and its Lin-  
 166 ear Programming (LP) relaxation. Across all configurations,  
 167 both algorithms achieved 100% feasibility and stability,  
 168 demonstrating their robustness at scale. Notably, DTAS  
 169 consistently matched the ILP-optimal objective, differing by  
 170 at most one or two app installations in rare instances. This  
 171 near-optimal performance was accompanied by impressive  
 172 computational efficiency, DTAS achieved near-optimal allo-  
 173 cations in few seconds, over 99% faster than the LP baseline,

175 while maintaining up to 30% runtime improvement over  
 176 CARL at scale. These results emphasize the practical scal-  
 177 ability and precision of our approach, positioning DTAS as  
 178 an effective balance between optimality and efficiency in  
 179 large-scale transaction allocation environments.

- 179 (4) **Dataset Generation.** We have used a semi-synthetic dataset  
 180 of transactions. Rabobank, a Dutch bank, provides data in the  
 181 form of cumulative transactions, and based on the reported  
 182 statistics of C2C (customer-to-customer), C2B (customer-to-  
 183 business), and B2B (business-to-business) transactions, this  
 184 dataset has been generated [11]. The data is presumed to be  
 185 from June 2023, and six years of transaction data have been  
 186 uniformly divided across 30 days of the month to simulate  
 187 daily activity. The dataset consists of four columns: *start id* (the ID of the user initiating the transaction), *value* (the  
 188 amount of money transferred), *day* (the day of the month  
 189 when the transaction occurred), and *end id* (the ID of the  
 190 user receiving the money).

190 Additionally, the distribution of UPI applications among  
 191 users has been generated using official NPCI statistics, which  
 192 provide the monthly transaction volumes handled by each  
 193 UPI app [8]. Each user in the dataset is assigned a set of  
 194 preinstalled UPI apps based on this real-world transaction  
 195 volume distribution. Specifically, higher-volume apps such as  
 196 PhonePe and GPay are more likely to be allocated to a larger  
 197 portion of users, while smaller-volume apps (e.g., Axis Bank,  
 198 Airtel Payments Bank) are assigned less frequently. This  
 199 probabilistic allocation reflects the real-world penetration  
 200 and usage patterns of UPI applications across the user base.

## 2 Background and Related Work

201 The Unified Payments Interface (UPI), developed by the National  
 202 Payments Corporation of India (NPCI) under the guidance of the  
 203 Reserve Bank of India (RBI), has emerged as a cornerstone of India's  
 204 digital economy. It serves as an interoperable architecture framework  
 205 equipped with standardized Application Programming Interfaces  
 206 (APIs) that enable seamless, real-time, and secure fund transfers  
 207 between individuals and merchants. UPI's open and inclusive design  
 208 allows multiple banks and fintechs to integrate easily, promoting  
 209 innovation and financial inclusion at scale [10].

210 As of 2024, UPI has over 350 million active users, connects  
 211 more than 550 banks, and powers 77 payment apps such as Google  
 212 Pay, PhonePe, BHIM, and WhatsApp Pay [1]. It has enabled the  
 213 deployment of over 340 million merchant QR codes across India. In  
 214 2023, UPI processed 117 billion transactions worth USD 2.19 trillion,  
 215 with Person-to-Merchant (P2M) payments accounting for around 62%  
 216 of total transactions [1]. UPI's phenomenal growth reflects India's  
 217 effective public-private partnership model, combining government  
 218 vision, RBI's progressive regulatory framework, and collaboration  
 219 among banks, fintechs, and merchants. However, rising transaction  
 220 volumes have simultaneously increased the risk of fraud and anomaly.  
 221 Several scholarly works, such as [2, 3, 5], have emphasized AI-driven  
 222 approaches to strengthen fraud detection mechanisms in digital  
 223 payment systems.

224 Beyond fraud detection, another critical concern in the financial  
 225 domain is fairness in decision-making. Unfair decision-making in

financial services occurs when certain groups or individuals face biased treatment in areas such as loan approvals, credit scoring, or mortgage access. Such bias can arise from historical discrimination or from machine learning models that inadvertently learn unfair patterns from data. Consequently, minority groups or equally qualified applicants may be denied fair financial opportunities, reinforcing existing social and economic inequalities. Song et al. [13] address this issue by proposing a Temporal Fair Graph Neural Network (TF-GNN) framework that models financial transactions as dynamic networks and enforces individual fairness over time. They introduce two new fairness notions specific to temporal graphs, provide a theoretical analysis of their fairness regret, and demonstrate through real-world experiments that their approach improves both prediction accuracy and fairness compared to prior methods.

There is another growing problem — the UPI ecosystem has become too concentrated, with over 80% of the 19.63 billion transactions (worth Rs. 24.90 lakh crore) processed in September 2025 being handled by just two third-party apps [7]. This duopoly poses serious systemic and security risks, as any disruption in these platforms could cripple a major share of India's digital payments. Moreover, such extreme concentration stifles innovation, reduces competition, and threatens fair market access within the country's largest payment infrastructure. To the best of our knowledge, this is the first work that formally studies and models this concentration issue, proposing mechanisms to enhance resilience and promote a more balanced and inclusive UPI ecosystem.

### 3 Problem Formulation

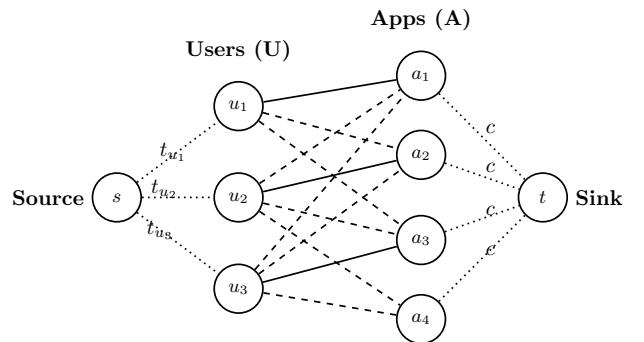
We model the problem of fairly balancing UPI transactions as a bipartite flow network. Let  $U$  denote the set of users and  $A$  denote the set of UPI apps. Each user  $u \in U$  generates  $t_u$  transactions, while each app  $a \in A$  can process at most  $c$  transactions, typically defined as a fraction of the total transaction volume. Users may already have certain apps installed, represented by solid edges  $E_{\text{solid}} \subseteq U \times A$ , while potential additional installations are represented by dashed edges  $E_{\text{dashed}} \subseteq U \times A$ . A source node  $s$  is connected to all users with edges of capacity  $t_u$ , and a sink node  $t$  is connected to all apps with edges of capacity  $c$ . The goal is to identify the minimal subset of dashed edges  $E' \subseteq E_{\text{dashed}}$  that need to be activated such that all transactions can be routed from  $s$  to  $t$  without exceeding app capacities or dropping transactions. Figure 1 illustrates our construction.

#### MINIMUM EDGE ACTIVATION FLOW (MEAF)

**Input.** Given a bipartite graph  $G = (U \cup A, E_{\text{solid}} \cup E_{\text{dashed}})$ , a source node  $s$  connects to each  $u \in U$  by an edge of capacity  $t_u$ , and a sink node  $t$  connects to each  $a \in A$  by an edge of capacity  $c$ . Edges in  $E_{\text{solid}}$  represent existing connections, while  $E_{\text{dashed}}$  denote optional edges that can be activated.

**Output:** Find the smallest subset  $E' \subseteq E_{\text{dashed}}$  such that all demands  $t_u$  can be routed from  $s$  to  $t$  through  $E_{\text{solid}} \cup E'$  without violating any capacity  $c$ .

Formally, let  $f(u, a)$  denote the number of transactions routed from user  $u$  to app  $a$ . The flow must satisfy the following conditions: every user's transactions are fully routed, i.e.,  $\sum_{a \in A} f(u, a) = t_u$ , and the total flow into any app respects its capacity,  $\sum_{u \in U} f(u, a) \leq c$ .



**Figure 1: Illustration of the bipartite UPI transaction flow network. Solid edges represent pre-installed apps ( $E_{\text{solid}}$ ), dashed edges denote potential installations ( $E_{\text{dashed}}$ ), and dotted edges indicate user transactions ( $t_{u_i}$ ) and app capacities ( $c$ ).**

Flow is allowed only on solid edges or activated dashed edges, and transaction units are indivisible, so  $f(u, a) \in \mathbb{Z}_{\geq 0}$ . Finally, the total flow in the network equals the total number of transactions, ensuring no transaction is dropped. The objective is to minimize  $|E'|$ , the number of additional app installations required to achieve a feasible integral flow. We now formally show that MINIMUM EDGE ACTIVATION FLOW is NP-complete.

**THEOREM 3.1.** MINIMUM EDGE ACTIVATION FLOW is NP-complete.

**PROOF:** It is easy to see that the problem is in NP: given a set of activated dashed edges and an integral flow assignment, we can verify in polynomial time that (i) flow conservation holds, (ii) capacities on  $(s, u)$  and  $(a, t)$  are respected, (iii) flow is sent only on activated dashed edges, and (iv) the number of activated dashed edges is at most  $k$ .

We prove NP-hardness by a polynomial-time reduction from 3-PARTITION. Let an instance of 3-PARTITION be given by  $S = \{s_1, \dots, s_{3m}\}$  and  $B$  with  $\sum_i s_i = mB$  and  $B/4 < s_i < B/2$  for all  $i$ . We construct an instance of UPI FLOW PROBLEM as follows.

**Construction.** Create  $3m$  users  $U = \{u_1, \dots, u_{3m}\}$  with  $t_{u_i} = s_i$ ,  $\forall i \in [3m]$ . Create  $m$  apps  $A = \{a_1, \dots, a_m\}$  with capacities  $c_{a_j} = B$  for all  $j$ . Add source  $s$  and sink  $t$  with edges  $(s, u_i)$  of capacity  $t_{u_i}$  and edges  $(a_j, t)$  of capacity  $B$ . Let  $E_{\text{solid}} = \emptyset$  and  $E_{\text{dashed}} = U \times A$  (i.e., every user can connect to every app via a dashed edge). All  $(u, a)$  edges have sufficiently large capacity (e.g., capacity  $t_u$ ) so that the app capacities are the binding constraints. Set the activation budget  $k := 3m$ .

This construction is clearly polynomial in the input size.

$(\Rightarrow)$  If the 3-PARTITION instance is a YES-instance, then the constructed flow instance admits a feasible integral flow using at most  $k = 3m$  activated dashed edges. Assume  $s$  can be partitioned into  $m$  disjoint triples  $T_1, \dots, T_m$  with  $\sum_{s \in T_j} s = B$  for each  $j$ . For each triple  $T_j$ , choose an app  $a_j$  and for each user  $u_i$  with  $s_i \in T_j$  activate exactly the dashed edge  $(u_i, a_j)$  and send the full amount  $t_{u_i} = s_i$  on that edge. For each app  $a_j$ , the incoming flow is  $\sum_{u_i \in T_j} t_{u_i} = B$ , so the capacity  $(a_j, t)$  is respected. Each user uses exactly one dashed edge, so the number of activated dashed edges is exactly  $3m = k$ . All source capacities  $(s, u_i)$  are saturated, hence the total flow value is

349  $\sum_i s_i = mB$ . Thus there exists a feasible integral flow using at most  $k$   
 350 activations.

351 ( $\Leftarrow$ ) If the constructed flow instance admits a feasible integral flow  
 352 using at most  $k = 3m$  activated dashed edges, then the 3-PARTITION  
 353 instance is a YES-instance. Suppose there is a feasible integral  $s-t$   
 354 flow of value  $\sum_{i=1}^{3m} s_i = mB$  using at most  $3m$  activated dashed edges.  
 355 Since  $E_{\text{solid}} = \emptyset$  and every user  $u_i$  has positive demand  $t_{u_i} = s_i > 0$ ,  
 356 each user must have at least one activated outgoing dashed edge in  
 357 order to send any flow. Therefore any feasible solution uses at least  
 358  $3m$  activated dashed edges. By the budget bound, the solution uses  
 359 exactly  $3m$  activations, hence each user activates *exactly one* dashed  
 360 edge and sends all of its (integral) demand through that edge.

361 Let  $S_j \subseteq U$  be the set of users assigned to app  $a_j$  (i.e., those for  
 362 which  $(u, a_j)$  is activated and carries the full  $t_u$ ). Flow feasibility  
 363 and capacity imply for every  $j$  that

$$\sum_{u \in S_j} t_u \leq c_{a_j} = B.$$

367 Summing over all apps and using that the total flow equals  $mB$  gives

$$\sum_{j=1}^m \sum_{u \in S_j} t_u = \sum_{u \in U} t_u = mB = \sum_{j=1}^m B,$$

371 which forces equality in each app separately:  $\sum_{u \in S_j} t_u = B$  for all  
 372  $j$ . Finally, by the 3-PARTITION bounds  $B/4 < t_u < B/2$ , no app  
 373 can receive 1 item (any  $t_u < B/2$ ) or  $\geq 4$  items (each  $t_u > B/4$   
 374 would exceed  $B$ ). Therefore each  $S_j$  has exactly three users and their  
 375 demands sum to  $B$ . The family  $\{S_1, \dots, S_m\}$  thus yields a partition  
 376 of  $s$  into  $m$  triples each summing to  $B$ , i.e., a YES-solution for  
 377 3-PARTITION.  $\square$

379 **Remark 3.2.** The MINIMUM EDGE ACTIVATION FLOW problem generalizes the classical SET COVER problem. Consequently, unless  
 380  $P = NP$ , it admits no polynomial-time approximation algorithm with  
 381 a factor better than  $(1 - o(1)) \log n$ .

384 This problem can be formulated as an Integer Linear Programming  
 385 (ILP) problem. We introduce variables  $f(u, a) \in \mathbb{Z}_{\geq 0}$  representing  
 386 the number of transactions routed from  $u$  to  $a$ , and binary variables  
 387  $x(u, a)$  for  $(u, a) \in E_{\text{dashed}}$  indicating whether a potential edge is  
 388 activated. The objective is to minimize the number of additional  
 389 edges activated:

$$\min \sum_{(u,a) \in E_{\text{dashed}}} x(u, a). \quad (1)$$

393 Each user's transactions must be fully routed, captured by

$$\sum_{a \in A} f(u, a) = t_u, \quad \forall u \in U, \quad (2)$$

396 while the capacities of apps must not be exceeded:

$$\sum_{u \in U} f(u, a) \leq c_a, \quad \forall a \in A. \quad (3)$$

400 Flow through a dashed edge is permitted only if the edge is activated:

$$f(u, a) \leq t_u \cdot x(u, a), \quad \forall (u, a) \in E_{\text{dashed}}. \quad (4)$$

403 Flow on solid edges is unconstrained by activation, i.e.,  $f(u, a) \geq 0$   
 404 for  $(u, a) \in E_{\text{solid}}$ . All flows are integral:

$$f(u, a) \in \mathbb{Z}_{\geq 0}, \quad x(u, a) \in \{0, 1\}. \quad (5)$$

## 4 Heuristic Algorithms for Minimum Edge Activation Flow

Given the computational hardness of the MINIMUM EDGE ACTIVATION FLOW problem, we first analyze a relaxed version of the ILP formulation to assess the strength of the model and the impact of relaxation on fairness and runtime. Specifically, the edge activation constraint in the ILP was relaxed such that activation variables corresponding to the dashed edges were allowed to take continuous values in the range  $[0, 1]$ , while retaining integrality for other decision variables. This relaxation is consistent with the modeling assumption that flow on solid edges is unconstrained by activation, i.e.,

$$0 \leq x(u, a) \leq 1 \quad \forall (u, a) \in E_{\text{dashed}}. \quad (6)$$

The relaxed ILP significantly reduced runtime, enabling faster convergence for larger data instances. However, despite this computational advantage, the integrality gap between the relaxed and original ILP solutions was observed to be nearly 0%, indicating that the ILP formulation is structurally tight and highly expressive. Furthermore, the relaxation did not yield notable improvements in fairness: the skewness in transaction distribution across applications remained largely unchanged.

These findings highlight that while the ILP formulation is strong, fairness and scalability improvements require algorithmic intervention beyond relaxation. To this end, we develop a series of **capacity-aware heuristic algorithms** that progressively enhance allocation efficiency, minimize redundant app installations, and maintain user trust stability. We begin with a baseline version of CARL, refine it based on empirical insights from transaction data, and finally introduce DTAS, which employs a two-stage allocation structure with an improved sorting logic to further enhance fairness and scalability.

### 4.1 Preliminary Capacity-Aware Reuse-First Layered Allocation

The initial version of CAPACITY-AWARE REUSE-FIRST LAYERED ALLOCATION (CARL) was developed to allocate transactions in a structured, layered manner with the objective of minimizing redundant app installations while ensuring full transaction coverage. In this preliminary design, users were processed in **descending order of their transaction-to-capacity ratio**, based on the intuition that heavy users should be prioritized to prevent app capacity overflow in later stages.

However, when tested on large-scale transaction datasets, this strategy exhibited significant inefficiencies due to the **highly skewed transaction distribution**. As illustrated in Figure 2, a small fraction of users generated the majority of transactions, while most users contributed only minimally. Processing heavy users first led to the early saturation of high-capacity applications, leaving lightweight users with limited or no access to their preinstalled apps. This forced many small users to install additional apps even when aggregate capacity was sufficient.

These effects manifested in two major shortcomings:

- (1) **Redundant installations:** Several lightweight users were compelled to install extra apps simply because their preferred high-capacity apps were prematurely exhausted.
- (2) **Poor utilization balance:** Heavy users dominated the capacity of major apps, while smaller apps remained underutilized.

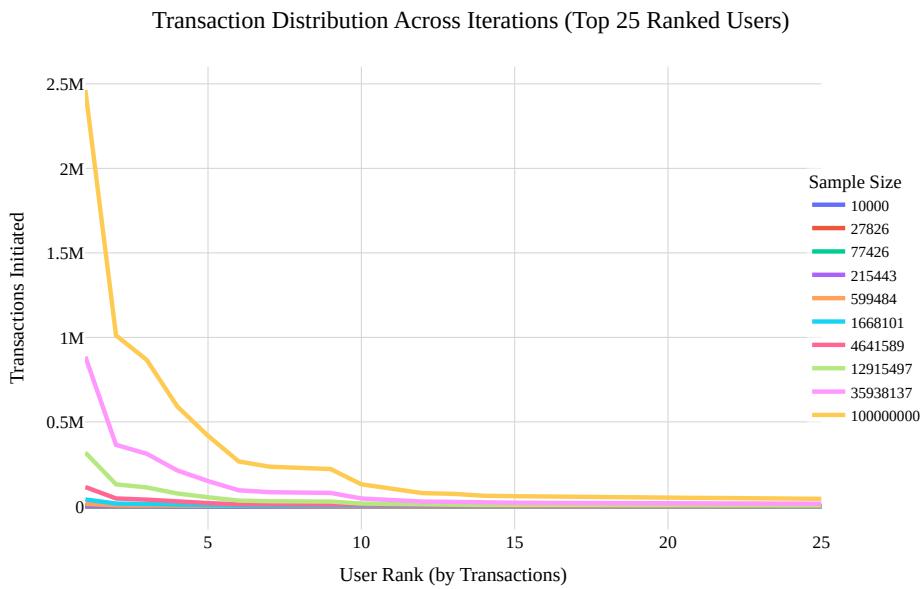


Figure 2: A few users generate most transactions, while the majority contribute very little.

Although Preliminary CARL achieved full transaction routing, it did so with suboptimal fairness, higher installation overhead, and reduced long-term scalability.

## 4.2 Refined CAPACITY-AWARE REUSE-FIRST LAYERED ALLOCATION

To address these shortcomings, the refined version of CARL reorders users in **ascending order** of their transaction-to-capacity ratio, effectively prioritizing small-volume users. This adjustment ensures that lightweight users benefit from available capacity early, resulting in better reuse of existing app capacity and a substantial reduction in redundant installations.

The refined algorithm assigns transactions through the following three allocation layers:

- (1) **Preinstalled Apps:** Users first utilize the apps they already have installed, incurring no additional installation cost.
- (2) **Existing Extra Apps:** If preinstalled apps lack sufficient capacity, the algorithm uses extra apps previously installed by any user.
- (3) **New Extra Apps:** Only as a last resort does the algorithm install new apps for a user, minimizing overall installation count.

This refined ordering yielded significant performance improvements. Under the experimental setup, the runtime dropped by nearly **99%**—from approximately 10,000 seconds to just **15 seconds**—while still producing results **very close to the optimal ILP** in terms of installation count and fairness.

Moreover, this *reuse-first layered allocation* not only enhances capacity utilization but also aligns with real-world behavioral constraints—users are reluctant to install multiple UPI apps, and incentive programs are often limited by budget. By recommending already adopted apps before introducing new ones, CARL fosters user trust,

reduces friction, and achieves a balanced transaction distribution with minimal intervention.

## 4.3 DECOUPLED TWO-STAGE ALLOCATION STRATEGY (DTAS)

Building on the insights gained from the refined CARL algorithm, we introduce the DTASf(DTAS), a significantly improved heuristic designed to minimize inter-user interference, prevent premature app saturation, and further reduce redundant installations. While CARL operates in a per-user, layer-by-layer fashion, DTAS fundamentally restructures the allocation workflow by **decoupling the allocation into two global phases**. This structural redesign eliminates several systemic inefficiencies observed in CARL and produces allocations that are consistently closer to the ILP-optimal solution.

**Two-Stage Allocation Framework.** Unlike CARL, which attempts to satisfy each user sequentially across multiple layers, DTAS processes all users collectively in two distinct stages:

- (1) **Stage 1: Preinstalled App Allocation.** In this stage, all users are allocated as many transactions as possible using only their preinstalled apps. By prioritizing preinstalled capacity upfront, DTAS ensures that every user's inherent, zero-cost connections are fully leveraged before the system begins introducing or reusing extra apps. This prevents later phases from unintentionally consuming the preinstalled capacity of users who rely on a limited number of apps.
- (2) **Stage 2: Extra App Allocation.** Once all preinstalled allocations are complete, the algorithm addresses the remaining unmet demand for each user. It first attempts to route additional transactions through **existing extra apps** that were previously installed for any user, promoting global reuse. Only if necessary—i.e., when no sufficient existing capacity

---

**Algorithm 1** CAPACITY-AWARE REUSE-FIRST LAYERED ALLOCATION

```

581 1: Initialize  $apps \leftarrow [0 \dots |initialApps| - 1]$ ,  

582    $remainingCapacity[app] \leftarrow c$   

583 2:  $transactionsHandled[app] \leftarrow 0$ ,  $extraApps \leftarrow \emptyset$ ,  

584    $userApps[u] \leftarrow mp[u]$   

585 3: for  $(u, apps)$  in  $mp$  do  $extraApps \leftarrow extraApps \cup apps$   

586 4: end for  

587 5:  $users \leftarrow [(id, mp[id], userTransactions[h])]$  for valid  $(id, h)$ ;  

588   Sort by  $\frac{|transactions|}{|preinstalled| \cdot c}$  desc  

589 6:  $totalRemaining \leftarrow 0$   

590 7: for  $(user, preinstalled, trans)$  in  $users$  do  

591 8:    $remaining \leftarrow |trans|$ ; sort  $preinstalled$  by  

592      $remainingCapacity$  desc  

593 9:   for  $a$  in  $preinstalled$  do  

594     if  $remaining \leq 0$  then break  

595   10:    end if  

596   11:     $alloc \leftarrow \min(remaining, remainingCapacity[a])$   

597   12:    if  $alloc > 0$  then  $remainingCapacity[a] -= alloc$ ;  

598      $transactionsHandled[a] += alloc$ ;  $remaining -= alloc$   

599   13:    end if  

600 14:   end for  

601 15:   if  $remaining > 0$  then  

602      $extras \leftarrow \{a \in extraApps \mid a \notin preinstalled\}$  sorted  

603       by  $remainingCapacity$   

604     for  $a$  in  $extras$  do  

605       if  $remaining \leq 0$  then break  

606     19:      end if  

607     20:       $alloc \leftarrow \min(remaining, remainingCapacity[a])$   

608     21:      if  $alloc > 0$  then  $remainingCapacity[a] -= alloc$ ;  

609        $transactionsHandled[a] += alloc$ ;  $remaining -= alloc$ ;  

610        $userApps[user] \cup= \{a\}$   

611     22:      end if  

612   23:    end for  

613 24:   end if  

614 25:   if  $remaining > 0$  then  

615      $newApps \leftarrow \{a \in apps \mid a \notin (preinstalled \cup$   

616        $extraApps)\}$  sorted by  $remainingCapacity$   

617     for  $a$  in  $newApps$  do  

618       if  $remaining \leq 0$  then break  

619     29:      end if  

620     30:       $alloc \leftarrow \min(remaining, remainingCapacity[a])$   

621     31:      if  $alloc > 0$  then  $remainingCapacity[a] -= alloc$ ;  

622        $transactionsHandled[a] += alloc$ ;  $remaining -= alloc$ ;  

623        $extraApps \cup= \{a\}$ ;  $userApps[user] \cup= \{a\}$   

624     32:      end if  

625   33:    end for  

626 34:   end if  

627 35:   end if  

628 36:    $totalRemaining += remaining$   

629 37: end for  

630 38:  $newMP \leftarrow \{u : list(userApps[u])\}$ ;  $extraAppsInstalled \leftarrow$   

631    $\sum_u |newMP[u]| - totalInitial$   

632 39: return  $(newMP, transactionsHandled, totalRemaining)$ 
```

---

is available—does DTAS install a **new extra app** for that user, thereby minimizing installation overhead.

This global decoupling is crucial. In CARL, later users could inadvertently saturate an app that served as the only preinstalled option for some earlier user, forcing the latter to install a new app unnecessarily. With DTAS, such conflicts are eliminated because all preinstalled utilizations are resolved before any extra-app allocations begin.

**Refined User Ordering.** To further improve fairness and capacity preservation, DTAS introduces a refined user-sorting mechanism. Users are processed in **ascending order of their transaction counts**. This prioritizes smaller users, ensuring they receive early access to available capacity.

For users with equal transaction volume, DTAS applies a second-level tie-breaker based on the **number of preinstalled apps**. Users with fewer preinstalled apps are routed earlier because they have fewer natural allocation options. This ordering prevents scenarios in which:

- ▷ a user with only one or two preinstalled apps loses access to those apps due to saturation by a more flexible user, or
- ▷ capacity is consumed by users who could have used alternate apps, leaving constrained users underserved.

This refined sorting ensures a more equitable distribution of limited capacity and preserves user trust by reducing unnecessary app installations.

**Advantages of DTAS Over CARL.** Through its two-stage separation and refined ordering, DTAS resolves key limitations observed in CARL:

- ▷ **No premature saturation of preinstalled apps.** Users retain access to their inherent app connections without being overridden by others.
- ▷ **Lower redundant installations.** DTAS consistently installs fewer extra apps across all datasets compared to both preliminary and refined CARL.
- ▷ **Improved stability and fairness.** Transaction loads are distributed more evenly across applications, preventing over-reliance on a small set of popular apps.
- ▷ **Closer to ILP-optimal results.** Empirical evaluations show that DTAS nearly matches the ILP's installation count, often differing by only one or two installations in large-scale experiments.

**Overall Impact.** By explicitly decoupling preinstalled and extra-app allocations and employing an ordering strategy aligned with fairness and capacity constraints, DTAS achieves a more stable, trust-preserving, and capacity-efficient transaction allocation. It builds on the strengths of CARL while eliminating its key limitations, resulting in a heuristic that is highly scalable, near-optimal in performance, and well suited for large-scale UPI transaction networks.

## 5 Experimental Evaluations

In this section, we present our empirical experiments and discuss the resulting performance insights.

639  
640  
641  
642  
643  
644  
645  
646  
647  
648  
649  
650  
651  
652  
653  
654  
655  
656  
657  
658  
659  
660  
661  
662  
663  
664  
665  
666  
667  
668  
669  
670  
671  
672  
673  
674  
675  
676  
677  
678  
679  
680  
681  
682  
683  
684  
685  
686  
687  
688  
689  
690  
691  
692  
693  
694  
695  
696

---

**Algorithm 2** DECOUPLED TWO-STAGE ALLOCATION STRATEGY (DTAS)

---

```

697 1: Initialize  $apps$ , set  $remainingCapacity[a] \leftarrow c$ ,  $handled[a] \leftarrow$ 
698   0
699 2:  $extraApps \leftarrow \bigcup_{u \in U} preinstalled[u]$  ▷ Apps preinstalled by
700   some user
701 3:  $userApps[u] \leftarrow preinstalled[u]$ 
702 4: Build  $users \leftarrow (u, preinstalled[u], transactions[u])$  and sort
703   by  $(|transactions|, |preinstalled|)$ 
704 5: Phase 1: Preinstalled Allocation
705 6: for  $(u, pre, txns)$  in  $users$  do
706   7:    $remain \leftarrow |txns|$ 
707   8:   for  $a$  in pre sorted by  $remainingCapacity[a]$  desc do
708     9:      $alloc \leftarrow \min(remain, remainingCapacity[a])$ 
709     10:     $remainingCapacity[a] -= alloc$ ;  $handled[a] += alloc$ ;
710        $remain -= alloc$ 
711     11:     if  $remain \leq 0$  then break
712     12:     end if
713   13:   end for
714   14:    $userRemain[u] \leftarrow remain$ 
715   end for
716 16: Phase 2: Extra App Allocation
717 17: for  $(u, pre, txns)$  in  $users$  do
718   18:    $remain \leftarrow userRemain[u]$ 
719   19:   if  $remain \leq 0$  then continue
720   20:   end if
721   21:   (a) Reuse existing extra apps
722   22:     for  $a$  in  $\text{sort}(extraApps \setminus pre, \text{by } remainingCapacity[a]$ 
723     desc}) do
724     22:      $alloc \leftarrow \min(remain, remainingCapacity[a])$ 
725     23:     Update  $remainingCapacity[a]$ ,  $handled[a]$ ,  $remain$ ,
726        $userApps[u]$ 
727     24:     if  $remain \leq 0$  then break
728     25:     end if
729   26:   end for
730   27:   (b) Install new extra app if needed
731   28:   if  $remain > 0$  then
732     28:      $a^* \leftarrow \arg \max_{a \notin pre \cup extraApps} remainingCapacity[a]$ 
733     29:     Allocate to  $a^*$ ; update  $remainingCapacity[a^*]$ ,
734        $handled[a^*]$ ,  $remain$ ,  $userApps[u]$ ,  $extraApps$ 
735     30:   end if
736   31: end for
737   32: Compute total extra installations and unallocated transactions
738   33: Return: final allocations and handled transactions

```

---

## 5.1 Computational Environment

The experiments were conducted on a Linux-based system with the following specifications: *Intel(R) Xeon(R) CPU E5-2630 v2 @ 2.60 GHz*, featuring *12 physical cores and 24 threads* across two sockets, with a total of *30 MB L3 cache* and two NUMA nodes. The system was equipped with *125 GiB of RAM* and *122 GiB of swap space*, and did not employ a dedicated GPU. It ran Ubuntu 22.04 LTS (64-bit). All algorithms were implemented in Python 3.12.3 and executed using the Gurobi Optimizer (version 12.0.3) as the solver for ILP formulations. Each configuration represents a unique combination of users, their respective transaction loads, and the available

UPI applications with varying capacity constraints. The ILP-based solution provides the optimal allocation and serves as the baseline for comparison. Heuristic approaches namely CAPACITY-AWARE REUSE-FIRST LAYERED ALLOCATION, and DECOUPLED TWO-STAGE ALLOCATION STRATEGY are evaluated in terms of their efficiency, scalability, and proximity to the optimal ILP solution.

In this section, we present our empirical experiments and discuss the resulting performance insights. Since UPI transaction data is not publicly available, we repurpose the Rabobank financial transaction dataset [12] as a proxy. The original multi-year dataset contains anonymized C2C, C2B, and B2B transfers; we preserve these proportions and uniformly redistribute the transactions over a 30-day window to obtain a realistic daily workload distribution. To simulate UPI behavior, we generate a population of users mapped to Rabobank account IDs and assign each user a set of preinstalled apps based on NPCI-reported monthly transaction shares. This probabilistic assignment ensures that high-volume apps (e.g., PhonePe, GPay) appear more frequently while still preserving a long tail of smaller apps, giving us a dataset that mirrors real-world UPI adoption patterns.

To assess the effectiveness and scalability of the proposed heuristics, we conducted a series of experiments comparing **ILP**, **LP Relaxation**, **CARL**, and **DTAS**. All algorithms were evaluated on synthetic configurations of increasing transaction volumes, ranging from 10,000 to 100 million transactions. Each configuration was tested under identical conditions to ensure fairness in comparison.

## 5.2 Performance Metrics

We evaluate each algorithm using the following metrics:

- ▷ **Number of App Installations:** Total number of app installations required to satisfy all transactions. Lower values indicate better efficiency.
- ▷ **Execution Time:** Total computational time taken to reach a feasible allocation.
- ▷ **Fairness Index:** A measure of how evenly transactions are distributed across apps (*Inverse Gini coefficient*).

**Comparison of App Installations:** Table 1 summarizes the total number of app installations required by each algorithm for increasing transaction sizes. The ILP provides the optimal lower bound, while the LP Relaxation closely approximates it. Among the heuristics, both CARL and DTAS achieve results close to the LP solution. Interestingly, CARL slightly outperforms DTAS for smaller instances due to its compact reuse-first ordering; however, as the dataset grows, DTAS becomes the superior choice—exhibiting better scalability, reduced redundant installations, and improved stability across large workloads.

**Execution Time Comparison:** Figure 3 presents the execution time of different algorithms across configurations. While ILP and LP approaches exhibit exponential growth in runtime with increasing data size, the heuristics (particularly DTAS) maintain near-linear scalability, making them viable for real-world deployment.

**Fairness Index:** As illustrated in Figure 4, the Gini index results were transformed into a fairness-oriented metric using the Inverse Gini Score (IGS = 1-G). DTAS achieves the highest fairness (IGS = 0.53), outperforming both ILP (0.34) and LP (0.30). This demonstrates DTAS's superior balance in capacity utilization and more equitable transaction distribution across UPI apps.

755  
756  
757  
758  
759  
760  
761  
762  
763  
764  
765  
766  
767  
768  
769  
770  
771  
772  
773  
774  
775  
776  
777  
778  
779  
780  
781  
782  
783  
784  
785  
786  
787  
788  
789  
790  
791  
792  
793  
794  
795  
796  
797  
798  
799  
800  
801  
802  
803  
804  
805  
806  
807  
808  
809  
810  
811

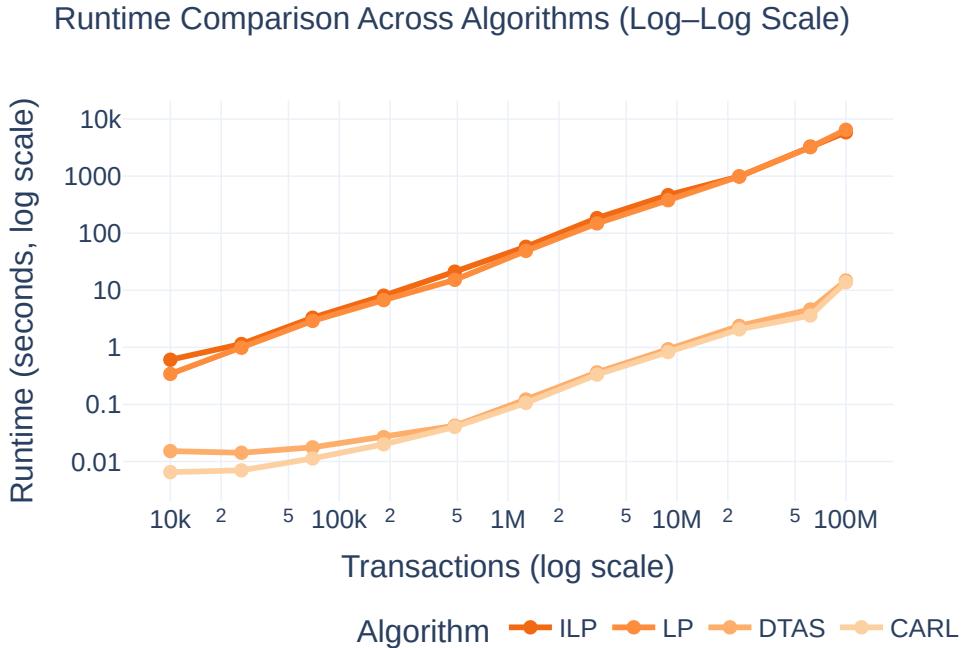


Figure 3: Execution time comparison of ILP, LP, CARL, and DTAS across increasing transaction volumes.

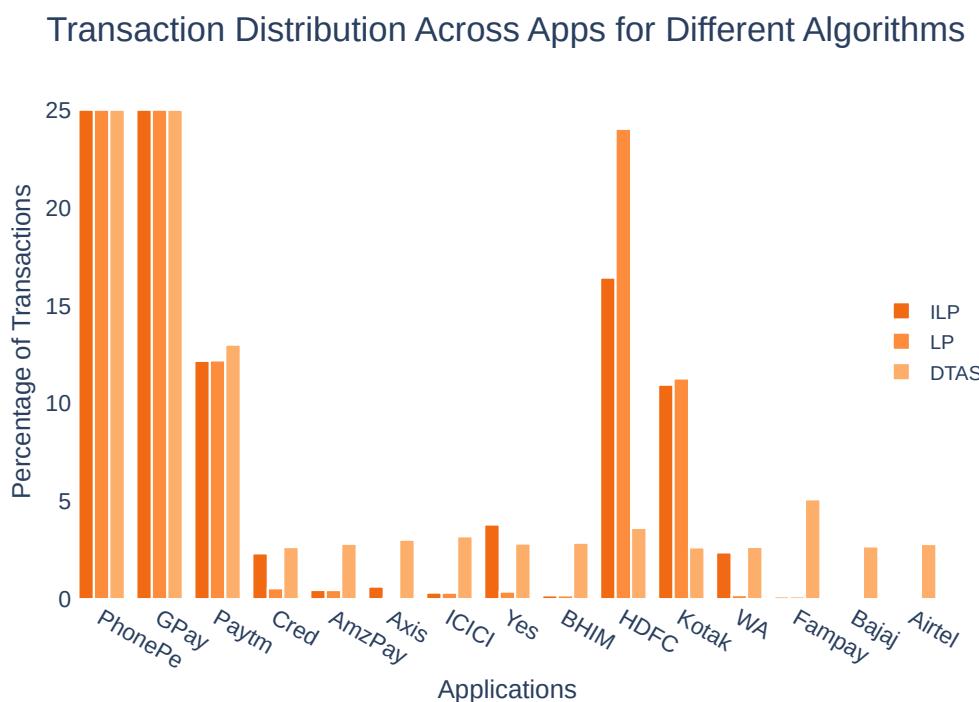


Figure 4: Transaction Allocation Across UPI Apps Using the Different Approaches on 25M Transactions.

Quantitative Results: The experiments were conducted on 1.21 million users and 100 million transactions to assess how varying

application capacity limits influence allocation outcomes. As shown in Table 2, increasing app capacity sharply reduces the number

**Table 1: Comparison of Installations Across Different Algorithms**

<b>Transactions</b>	<b>ILP</b>	<b>LP</b>	<b>CARL</b>	<b>DTAS</b>
10,000	5	3.78	5	5
69,519	27	26.03	29	28
483,293	90	88.99	95	91
1,274,274	261	260.32	273	261
8,858,667	1,901	1,900.40	1,967	1,902
23,357,214	4,697	4,696.06	4,855	4,697
100,000,000	***	***	6,863	6,599

Note: ‘\*\*\*’ indicates that the runtime for ILP and LP was prohibitively high, and results could not be computed within reasonable time limits.

of users requiring new installations—from about 2.25% at a 10% capacity limit to nearly zero at 40%. This implies that a vast majority of users can be fully served even under moderate capacity levels, and only a very small subset (less than 2.3%) faces unmet demand.

However, this increased capacity also accentuates duopoly tendencies, where a few dominant applications absorb a disproportionate share of the total transaction load. Such concentration, while improving feasibility, may negatively affect competition and resilience. Conversely, lowering app capacity reduces dominance but results in unallocated transactions—over 65 million at a 10% limit—indicating potential service degradation if strict caps are imposed. If regulatory constraints, such as NPCI’s transaction thresholds per app, were enforced, the corresponding remaining transactions directly represent the volume that would go unserved, potentially leading to user friction and reduced system efficiency.

**Table 2: Impact of App Capacity on Remaining Demand and Unallocated Transactions**

<b>App Cap.</b>	<b>Users w/ Remaining Demand (%)</b>	<b>Unallocated Transactions</b>
10	2.25%	65,982,913
15	1.00%	50,982,913
20	0.54%	40,418,301
25	0.29%	30,418,301
30	0.13%	20,418,301
35	0.02%	10,418,301
40	≈0%	1,525,685

## 6 Conclusion

The UPI ecosystem is heavily dominated by just two to three apps, raising concerns of duopoly. NPCI’s 30% cap aims to address this imbalance, but enforcing such a limit in practice turns out to be computationally challenging. In this work, we take a first step toward understanding this problem and propose DTAS, a simple and scalable strategy that keeps users’ existing app choices intact and adds the fewest new installations needed for compliance. Our experiments show that DTAS performs almost as well as the LP baseline while running efficiently even at very large scales.

Our study focuses on balancing transaction counts across apps. A natural next direction is to incorporate transaction value, since high-value transfers may strain capacity differently, influence user incentives, and expose apps to disproportionate financial risk. Extending the framework to value-aware allocation could therefore reveal

richer fairness constraints and more realistic resilience challenges for the UPI ecosystem.

## References

- [1] European Payments Council AISBL. 2024. UPI: revolutionising real-time digital payments in India. <https://www.europeanpaymentscouncil.eu/news-insights/insight/upi-revolutionising-real-time-digital-payments-india>. Accessed: 2024-06-26.
- [2] Deepa Devassy, Divya K S, Gipsy Paul Mannickathan, Alen Biju, Aswathi T, and Devarsh R. 2025. FedUPI: Federated Learning Empowered Detection of Fraudulent UPI Transactions. In *2025 5th International Conference on Soft Computing for Security Applications (ICSCSA)*. 1007–1015. doi:10.1109/ICSCSA66339.2025.11170880
- [3] Muhammad Liman Gambo, Anazida Zainal, and Mohamad Nizam Kassim. 2022. A convolutional neural network model for credit card fraud detection. In *2022 International Conference on Data Science and Its Applications (ICoDSA)*. IEEE, 198–202.
- [4] Jaspreet Kalra. 2024. India delays UPI market share cap in relief for Walmart backed PhonePe, Google Pay. [www.reuters.com/technology/india-delays-upi-payments-market-share-cap-relief-walmart-backed-phonepe-google-2024-12-31/](http://www.reuters.com/technology/india-delays-upi-payments-market-share-cap-relief-walmart-backed-phonepe-google-2024-12-31/).
- [5] N Lingareddy, Devadutta Indoria, S Deepika, Geogen George, M Keerthi Priya, et al. 2025. Enhancing Digital Payment Security: UPI Fraud Detection with Advanced Machine Learning Algorithms. In *2025 Global Conference in Emerging Technology (GINOTECH)*. IEEE, 1–7.
- [6] Amar A Mahawish and Hassan J Hassan. 2022. Improving RED algorithm congestion control by using the Markov decision process. *Scientific Reports* 12, 1 (2022).
- [7] Meghna Mittal. 2025. 80% of UPI volumes concentrated with two app providers, fintech body writes to FinMin. <https://shorturl.at/LFhTN>.
- [8] National Payments Corporation of India (NPCI). 2025. Unified Payments Interface (UPI) Ecosystem Statistics. <https://www.npci.org.in/what-we-do/upi/upi-ecosystem-statistics>.
- [9] PIB. 2025. India’s UPI Revolution. <https://www.pib.gov.in/PressNoteDetails.aspx?NoteId=154912&ModuleId=3>.
- [10] Shailesh Rastogi, Chetan Panse, Arpita Sharma, and Venkata Mrudula Bhimavarapu. 2021. Unified Payment Interface (UPI): A digital innovation and its impact on financial inclusion and economic development. *Universal Journal of Accounting and Finance* 9, 3 (2021), 518–530.
- [11] Akrati Saxena. 2023. RaboBank Transaction Dataset. [https://github.com/akratiet/RaboBank\\_Dataset](https://github.com/akratiet/RaboBank_Dataset). Accessed: 2025-04-12.
- [12] Akrati Saxena, Yulong Pei, Jan Veldsink, Werner van Ipenburg, George Fletcher, and Mykola Pechenizkiy. 2021. The banking transactions dataset and its comparative analysis with scale-free networks. In *Proc. ASONAM*. 283–296.
- [13] Zixing Song, Yuji Zhang, and Irwin King. 2023. Towards fair financial services for all: A temporal GNN approach for individual fairness on transaction networks. In *Proceedings of the 32nd ACM international conference on information and knowledge management*. 2331–2341.